Machine Learning Based Prediction of CR Severity Level in FLOSS: Experimental Results

Luiz Alberto Ferreira Gomes
gomes.luiz@ic.unicamp.br
Institute of Computing, University of Campinas, Brazil.

Mario Lucio Cortes
cortes@ic.unicamp.br
Institute of Computing, University of Campinas, Brazil

Abstract—In the context of Change Request (CR) systems, the severity level of a change request is considered a critical variable when planning software maintenance activities, indicating how soon a CR needs to be addressed. However, the severity level assignment remains primarily a manual process, mostly depending on the experience and expertise of the person who has reported the CR. This paper presents preliminary findings on the prediction of CR severity level by analyzing its long description, using text mining techniques and Machine Learning (ML) algorithms. We have collected CRs from three FLOSS projects (imbalanced) repositories: Cassandra, Hadoop and Spark. Ours results were better than those published in the literature in terms of F-measure performance for two research questions (using Random Forest) and similar for the third research question. However, subsequent analyses based on the Friedman test have demonstrated that data used in experiments haven't permitted us to say with enough confidence level that Random Forest is better than the others ML algorithms. We have also shown that the use classical ML measurements available in the literature may not help deciding whether a ML approach will bring any benefit to the user, and have proposed an alternative measuring approach to address this issue.

Index Terms—software maintenance; change request systems; text mining; machine learning; software repositories.

I. INTRODUCTION

Change Request (CR) systems have played a major role in maintenance process in many software development settings, both in Closed Source Software (CSS) and in Open Source Software (OSS) scenarios. This is especially true in OSS, which is characterized by the existence of many of users and developers with different levels of expertise spread out around the world, who might create or be responsible for dealing with several CRs [1].

A user interacts with a CR system often through a simple mechanism called CR form. This form enables him to request changes, to report bugs or to ask for support in a software product [2]. Initially, he or she should inform a short description, a long description, a type (e.g. bug, new feature, enhancement, and task) and an associated severity level (e.g. blocker, critical, major, minor and trivial). Subsequently, a development team member will review this request and, case it is not refused for some reason (e.g. request duplication), he or she will complete the information in CR form, indicating, for example, its priority and assigning the person responsible for the CR.

The severity level information is recognized as a critical variable in the equation to estimate a prioritization of CRs [3]. It defines how soon the CR needs to be addressed [4]. However, the severity level assignment remains mostly a manual process which relies only on the experience and expertise of the person who has opened the CR [1], [3], [4]. Consequently, it is a process with high degree of subjectivity, and it may be quite error-prone.

The number of CRs in large and medium software OSS projects is frequently very large [5]: Eclipse project received over 2.764 requests and GNOME project received over 3.263 requests between 01/10/2009 and 01/01/2010. Severity level shifts throughout CR lifecycle may have an adverse effect on planning of maintenance activities. For example, the maintenance team could be assigned to address less significant CRs before most important ones. There has been reports [6] of efforts to implement intelligent software assistants to help developers and maintenance personnel in defining more accurately the field values in a CR form. Currently, Machine Learning techniques have become a popular method to address this issue and there are quite a few publications in this area in the literature [1].

Machine Learning (ML) techniques have been successfully applied in solving real problems in many areas of knowledge, including those related to CR systems, such as duplication and assignment of CR [1]. However, the accuracy of ML algorithms may be affected by imbalanced datasets [7] —a recurring critical problem in CR repositories [6]. For example, more than 60% of CRs may have a "major" severity level. In addition to this problem, most publications are still focused in predicting the severity level of CRs and none of them have been implemented into popular tools like as Bugzilla, Jira and Redmine. [1]. Furthermore, many have used proprietary and/or not public ML algorithms. Therefore, there is still a clear need of advances in this knowledge area, specially broadening the reach of research questions and including more popular and open OSS and ML algorithms.

In this context, the general purpose of our research is to develop an intelligent ML based assistant to help developers and maintenance personnel in the OSS maintenance activities. In this current article, our specific goals are:

G_1. Evaluate the performance of traditional ML algorithms in the prediction of CR severity level and identify a suitable algorithm to perform such prediction in a scenario where imbalanced data is natural;

G$_2$. Apply statistical tests to confirm if a ML algorithm is better than the others.

G$_3$. Analyze whether ML algorithms outperforms a human user in predicting CR severity level and propose new approaches, accordingly.

To meet these goals, this research works with the following research questions, regarding CR severity level during its lifecycle:

RQ$_1$. *Will the CR severity level change?*

RQ$_2$. *Will the CR severity level increase, decrease or remain the same*?

RQ$_3$. *What is the prediction for the final CR severity level?*

RQ$_4$. *How ML predictions compare to user prediction?*

The first three research questions are related to goals 1 and 2, and the last one is related to goal 3

The contributions of our research are:

- Analyze the performance of three popular ML algorithms as multi category classifiers in imbalanced datasets extracted from three FLOSS repositories.
- Run statistical tests on experimental results and demonstrate that, under these conditions, it is not possible to guarantee that an algorithm is better than others with appropriate confidence level.
- Analyze how well an automated ML based predictor performs with respect to a user prediction and propose new measurement approach.

The article is organized as follows. Section II presents related work that are relevant to our research. Section III provides the information background about CR systems, text mining and machine learning techniques necessary to understand our approach. Section IV describes our work. Section V presents final findings and discussion. Finally, Section VI present conclusions and future work.

II. RELATED WORK

This section presents relevant articles in mining open system repositories, aiming at extracting data and using ML techniques to predict several maintenance properties.

Menzies and Marcus [8] have developed a method, named SEVERIS (SEVERity ISsue assessment), for evaluating the severity of CRs. SEVERIS is based on established data and text mining techniques. The method was applied to predict CR severity level in five projects managed by the Project and Issue Tracking System (PITS), an issue tracker system used by NASA (Stratified F-measures by severity level in the range: (2) 78%-86%; (3) 68%-98%; (4) 86%-92%).

Lamkanfi et al. [4] have developed an approach to predict if severity of bug report is non-severe (severity levels: 1 or 2) or severe (severity levels: 4 or 5) based on text mining algorithms (tokenization, stop word removal, stemming) and on the Naïve Bayes machine learning algorithm. They have validated their approach with data from three open source project (Mozilla, Eclipse, and GNOME). The article reports that a training set with approximately 500 CRs per severity level is sufficient to make predictions with reasonable accuracy (precision and recall in the range 0.65-0.75 with Mozilla and Eclipse; 0.70-0.85 with GNOME).

Valdivia et al. [9] have characterized blocking bugs in six open source projects and proposed a model to predict them. Their model was composed of 14 distinct factors or features (e.g. the textual description, location the bug is found in and the people involved with the bug). Based on these factors they have built decision trees for each project to predict whether a bug will be a blocking bug or not (F-measures in the range 15-42%).

Tian et al. [3] have develop a method to predict the severity level of new CRs based on similar CRs reported in the past. The comparison between old and new CRs was implemented by the BM25 similarity function. This method was applied to Mozilla, Eclipse and OpenOffice projects over more than 250,000 CR extracted from Bugzilla (F-measure in the range 13.9-65.3% for Mozilla; 8.6-58% for Eclipse; and 12.3-74% for OpenOffice).

III. BACKGROUND

This section briefly comments of basic concepts necessary to understand this research area, namely CR Systems, Text Mining, Machine Learning, and ML evaluation metrics.

Data used in this research area are usually extracted from the so-called CR Systems, or Bug Tracking Systems. Popular CR Systems are Bugzilla, Jira, and Redmine [3]. In this work, we extracted Cassandra, HADOOP and Spark datasets from Jira CR System. Additional information can be found in [10].

Two techniques are frequently used in this research area: Text Mining [11] [12] and Machine Learning (ML) [12] [13] [14] [15]. Detailing of these techniques are outside the scope of this paper.

Finally, it is worth mentioning the specific metrics we use for assessing prediction performance. The three most common performance measures for evaluating the accuracy of classification algorithms are precision, recall, and F-measure, described as follows [16] [17]:

Recall. Recall is the number of True Positives (TP) divided by the number of True Positives (TP) and of False Negatives (FN), where the TP and FN values are derived from the confusion matrix. A low recall indicates many false negatives.

Precision. Precision is the number of True Positives (TP) divided by the number of True Positives and False Positives (FP). A low precision can also indicate many false positives.

F-measure. F-measure conveys the balance between precision and recall, and can be calculated as their harmonic mean.

IV. EXPERIMENT

This section describes the experiment conducted to address the Research Questions. As in typical methodologies used in ML studies, it comprises the following steps: Data Collection (IV-A), Data Preprocessing (Section IV-B), and Training and Testing (Section IV-C).

A. Data Collection

This step in the experimental research encompasses selecting FLOSS datasets to serve as the data source, studying and interpreting its data structure, and finally extracting relevant data from its repository (feature extraction). In this research, Cassandra, Hadoop, Linux, Mozilla, and Spark Open Systems were considered as potential Open Source Systems to study. In a first approximation, Cassandra, Hadoop and Spark were selected as data sources of CR records, due to the fact they are open, well stablished, have a considerable number of CRs already registered, use standard repositories, and were under study by other researchers in our research group.

Cassandra[cassandra.apache.org] is a distributed NoSQL database management system designed to handle large amounts of data across many servers, providing fault-tolerance with no single point of failure. Hadoop[hadoop.apache.org] is a framework that allows for the distributed processing of large data sets across clusters of computers using simple programming models. Spark[spark.apache.org] is a cluster-computing engine for large-scale data processing which provides an interface for programming entire clusters with implicit data parallelism and fault-tolerance. They are considered a specialized and complex FLOSS project with many users with different levels of expertise.

CRs from these FLOSS projects are stored in a Jira based repository [https://www.atlassian.com/software/jira] which allows for access to all CR contents in XML format. Everything is available (except change history), from CR long description field (with lines with few characters to ones with many lines), including code snippets and exception stack trace. Two steps are used to perform data extraction from their web site[http://issues.apache.org]: (i) copying CR basic data (e.g. status and resolution) from XML contents; and (ii) copying CR changes history from external HTML pages (this may be important for learning).

CR record data from February 01, 2006 to May 07, 2017 were collected. The total number of CR records retrieved after preprocessing was 22901.

Figure 1 shows how the 7538 retrieved CR Cassandra records were distributed in terms of severity level and severity level change. Figure 1(a) shows the severity level distribution: 9.7% have severity trivial (1); 37.6% have severity minor (2); 48.4% have severity major (3), 3.0% have severity critical (4), and 1.3% have severity blocker (5). Figure 1(b) shows that only 7.0% have changed their severities levels during the CR lifecycle. Finally, Figure 1(c) reveals that of these 7.0% CRs which changed their severity, 67% decreased it, and 33% increased it.

Figure 2 shows how the 8262 retrieved CR Hadoop records from were distributed in terms of severity level and severity level change. Figure 2(a) shows the severity level distribution: 4.3% have severity trivial (1); 19.6% have severity minor (2); 61.2% have severity major (3), 3.8% have severity critical (4), and 11.1% have severity blocker (5). Figure 2(b) shows that only 8.0% have changed their severities levels during the CR

lifecycle. Finally, Figure 2(c) reveals that of these 8.0% CRs which changed their severity, 30.8% decreased it, and 69.2% increased it.

Figure 3 shows how the 7101 retrieved CR Spark records were distributed in terms of severity level and severity level change. Figure 3(a) shows the severity level distribution: 2.9% have severity trivial (1); 4.4% have severity minor (2); 22.5% have severity major (3), 50.6% have severity critical (4), and 10.1% have severity blocker (5). Figure 3(b) shows that only 13.3% have changed their severities levels during the CR lifecycle. Finally, Figure 3(c) reveals that of these 13.3% CRs which changed their severity, 35.9% decreased it, and 64.1% increased it.

Summarizing findings in Figures 1, 2 and 3: (a) the most frequent severity level type is "major"; (b) there have been few changes in severity levels; (c) in two of them (Hadoop and Spark), severity levels increased. Furthermore, one can see that these datasets are clearly imbalanced, posing additional difficulty to the application of the ML methodology.

B. Preprocessing

Raw data previously collected from the Cassandra, Hadoop and Spark CR repositories were not properly structured to serve as input to ML algorithms, they weren't in tidy data format [18]. The classical way to address this problem is to run preprocessing procedures to extract, organize and structure relevant features out of the raw data. Specific scripts were written in R language to accomplish this features extraction. Preprocessing tasks were executed as follows:

- Extraction of relevant features: key, type, status, resolution status, days to resolve, quantity of comments, severity level and long description of CRs;
- Selecting only CRs with status equals to Closed and resolution equals to Fixed and Implemented. This type of CRs was effectively implemented by the development team and they can no longer have their severity level changed.
- Merging CR features with their change history data. This additional information allows for the identification of CRs that have changed severity level during the CR lifecycle, and furthermore, if they have changed for better (decrease) or worse (increase).
- Performing text mining in the long description field to identify the 100 most frequent words. This information is then converted into features for each CR.

C. Training and testing

Training and testing steps start with partitioning the already preprocessed dataset in two disjoint subsets: a subset for training, with 80% of the CRs, and a subset for testing, with the remaining 20% of the CRs. Three classical sampling approaches, random, proportional, and uniform [19] were analyzed to select the training set. Best results were obtained with the random sampling technique. In the training phase, we have used the 5×3 Repeated Cross-Validation technique [17] to obtain more stable estimates of each algorithm's performance

Fig. 1: Cassandra dataset distributions: (a) severity level (b) change pattern (c) direction of change.

Fig. 2: Hadoop dataset distributions: (a) severity level (b) change pattern (c) direction of change.


Fig. 3: Spark dataset distributions: (a) severity level (b) change pattern (c) direction of change.

and enhance replicability of the results [19]. In the testing phase, each ML algorithm was validated with 20% of each CR dataset to measure its accuracy.

We have chosen three traditional ML algorithms: Neural Networks [20], Random Forest [15] and Support Vector Machine(SVM) [21] which were implemented, respectively, using neuralnet (with Single Hidden Layer), randomForest, and kernlab (with Radial Basis Function Kernel and multi-class classification) R libraries.

V. FINDINGS AND DISCUSSIONS

This section presents the experimental findings listed by research question: RQ1 (V-A), RQ2 (Section V-B), RQ3 (Section V-C) and RQ4 (Section V-D). In addition, the performance of ML algorithms is evaluated using Friedman Statistical Test (Section V-E).

A. RQ1: Will the CR severity level change?

The RQ1 is a simple binary problem, i.e., a question whose answer is true (class 1) or false (class 0). Tables I, II and III show the performance of ML algorithms to predict the response to this issue.

We tested the ML algorithms with 4580 (20% of 22901) CRs: 4154 have changed their severity level, and 426 haven't

TABLE I: ML algorithms precision performance on RQ1.

	Class	Neural Network	Random Forest	SVM
		Precision		
Cassandra	0	0.9530591	0.9596013	0.9581371
	1	0.7241379	0.9740260	1.0000000
Hadoop	0	0.9530516	0.9498208	0.9477157
	1	0.6339286	0.8289474	0.9830508
Spark	0	0.9125249	0.9274406	0.9134555
	1	0.6162162	0.7640449	0.9819820
	Average	0.7988197	0.9006468	0.9640569

changed their severity level. We can observe that the three algorithms performed very closely. However, the Random Forest algorithm had the best performance in two out of three measures.

We have also analyzed the percentage of incorrect predictions made in answering RQ1. Figure 4 shows that this implementation of the Neural Network Algorithm had the worst performance and the two others performed similarly.

Fig. 4: Performance of ML algorithms for RQ1 and RQ2.

TABLE II: ML algorithms recall performance on RQ1.

	Class	Neural Network	Random Forest	SVM
		Recall		
Cassandra	0	0.9868924	0.9989077	1.0000000
	1	0.4144737	0.4934211	0.4736842
Hadoop	0	0.9780514	0.9930407	0.9994647
	1	0.4409938	0.3913043	0.3602484
Spark	0	0.9509669	0.9709945	0.9986188
	1	0.4634146	0.5528455	0.4430894
	Average	0.7057988	0.7334190	0.7125176

TABLE III: ML algorithms F-measure performance on RQ1.

	Class	Neural Network	Random Forest	SVM
		F-measure		
Cassandra	0	0.9696807	0.9788600	0.9786211
	1	0.5271967	0.6550218	0.6428571
Hadoop	0	0.9653897	0.9709500	0.9729026
	1	0.5201465	0.5316456	0.5272727
Spark	0	0.9313493	0.9487179	0.9541405
	1	0.5290023	0.6415094	0.6106443
	Average	0.7404609	0.7877841	0.7810730

B. RQ2: Will the CR severity level increase, decrease or remain the same?

The RQ2 poses a problem more difficult than the previous question. It is a question with three possible responses related to severity level: it has decreased (class -1); it has remained (class 0); and it has increased (class 1). Tables IV, V, and VI shows the performance of the ML algorithms to predict the response to this issue.

TABLE IV: ML algorithms precision performance on RQ2.

	Class	Neural Network	Random Forest	SVM
		Precision		
Cassandra	-1	0.4393939	0.7435897	0.7631579
	0	0.9523305	0.9565445	0.9546403
	1	0.7333333	0.7428571	0.8214286
Hadoop	-1	0.2272727	0.5555556	0.5000000
	0	0.9549266	0.9551084	0.9520653
	1	0.6060606	0.7195122	0.8666667
Spark	-1	0.2500000	0.5925926	0.5769231
	0	0.9119788	0.9303548	0.9157695
	1	0.5555556	0.6689655	0.7977528
	Average	0.6256502	0.7627867	0.7942671

We tested the ML algorithms with 4580 (20% of 22901) CR. Only now, we have three predicting situations: 4154 haven't

TABLE V: ML algorithms recall performance on RQ2.

	Class	Neural Network	Random Forest	SVM
		Recall		
Cassandra	-1	0.2989691	0.2989691	0.29896907
Cassandra	0	0.9819771	0.9978154	1.00000000
Cassandra	1	0.3928571	0.4642857	0.41071429
Hadoop	-1	0.1111111	0.1111111	0.08888889
Hadoop	0	0.9753747	0.9908994	0.99946467
Hadoop	1	0.5172414	0.5086207	0.44827586
Spark	-1	0.1500000	0.2000000	0.18750000
Spark	0	0.9516575	0.9779006	0.99861878
Spark	1	0.4518072	0.5843373	0.42771084
	Average	0.5367772	0.5704376	0.5400158

TABLE VI: ML algorithms F-measure performance on RQ2.

	Class	Neural Network	Random Forest	SVM
		F-measure		
Cassandra	-1	0.3558282	0.4264706	0.4296296
Cassandra	0	0.9669266	0.9767442	0.9767938
Cassandra	1	0.5116279	0.5714286	0.5476190
Hadoop	-1	0.1492537	0.1851852	0.1509434
Hadoop	0	0.9650424	0.9726747	0.9751893
Hadoop	1	0.5581395	0.5959596	0.5909091
Spark	-1	0.1875000	0.2990654	0.2830189
Spark	0	0.9313957	0.9535354	0.5290023
Spark	1	0.4983389	0.6237942	0.5568627
	Average	0.5693392	0.6227619	0.6073741

TABLE VII: ML algorithms precision performance on RQ3.

	Class	Neural Network	Random Forest	SVM
		Precision		
Cassandra	1	0.4022989	0.6976744	0.5348837
Cassandra	2	0.5394737	0.6209440	0.7513966
Cassandra	3	0.6645221	0.6924959	0.6222510
Cassandra	4	0.4782609	1.0000000	1.0000000
Cassandra	5	0.6666667	1.0000000	1.0000000
Hadoop	1	0.2000000	0.9333333	0.8750000
Hadoop	2	0.3964497	0.7433628	0.9452055
Hadoop	3	0.6668558	0.7057175	0.6960305
Hadoop	4	0.3333333	1.0000000	1.0000000
Hadoop	5	0.4032258	0.9204545	1.0000000
Spark	1	0.1818182	0.7142857	0.8000000
Spark	2	0.3345070	0.4785276	0.8194444
Spark	3	0.6096892	0.6131657	0.5994532
Spark	4	0.4807692	0.9733333	0.9726027
Spark	5	0.3703704	0.7857143	0.9500000
	Average	0.4485493	0.7919339	0.8377511

TABLE VIII: ML algorithms recall performance on RQ3.

	Class	Neural Network	Random Forest	SVM
		Recall		
Cassandra	1	0.2243589	0.1923077	0.1474359
Cassandra	2	0.5766526	0.5921238	0.3783404
Cassandra	3	0.7053658	0.8282927	0.9385366
Cassandra	4	0.3283582	0.4626866	0.4626866
Cassandra	5	0.0800000	0.2400000	0.2400000
Hadoop	1	0.0131578	0.1842105	0.1842105
Hadoop	2	0.1850828	0.2320442	0.1906077
Hadoop	3	0.9136858	0.9790047	0.9953344
Hadoop	4	0.1265822	0.3544304	0.3417722
Hadoop	5	0.1111111	0.3600000	0.3244444
Spark	1	0.0312500	0.0781250	0.0625000
Spark	2	0.2691218	0.2209632	0.1671388
Spark	3	0.7494382	0.9314607	0.9853933
Spark	4	0.3989361	0.3882979	0.3776596
Spark	5	0.2531645	0.2784810	0.2405063
	Average	0.3310844	0.4214952	0.4024377

changed their severity level, 246 have increased their severity level, and 180 have decreased their severity level. We can observe in Tables IV, V and VI which the ML algorithms also performed very closely as question 1. However, the Random Forest algorithm had also the best performance in two out of three measures.

Figure 4 shows that the ML algorithms performance had the same pattern in RQ2, as compared to RQ1.

C. RQ3: What is the prediction for the final CR severity level?

The RQ3 is a problem much harder than other two. It is a question with five responses related to severity level: (1) trivial; (2) minor; (3) major; (4) critical; and (5) blocker. Tables VII, VII and VII shows the performance of the ML algorithms to predict the response to this issue.

We tested the ML algorithms with 4580 (20% of 22901) CRs. Only now, we have six predicting situations: 288 are trivial; 1218 are minor; 2470 are major; 259 are critical; 345 are a blocker. We can observe in the Table VII, VIII and IX which the ML algorithms also performed very closely as questions 1 and 2. As in the two previous questions, the

TABLE IX: ML algorithms F-measure performance on RQ3.

	Class	Neural Network	Random Forest	SVM
		F-measure		
Cassandra	1	0.2880658	0.3015075	0.2311558
	2	0.5574439	0.6061915	0.5032741
	3	0.6843350	0.7543314	0.7483469
	4	0.3893805	0.6326531	0.6326531
	5	0.1428571	0.3870968	0.3870968
Hadoop	1	0.0246913	0.3076923	0.3043478
	2	0.2523540	0.3536842	0.3172414
	3	0.7709973	0.8201954	0.8192000
	4	0.1834862	0.5233645	0.5094340
	5	0.1742160	0.5175719	0.4899329
Spark	1	0.0533333	0.1408451	0.1159420
	2	0.2982731	0.3023256	0.2776471
	3	0.6723790	0.7395183	0.7454314
	4	0.4360465	0.5551331	0.5440613
	5	0.3007518	0.4112150	0.3838384
	Average	0.3485740	0.4902217	0.4673068

Random Forest algorithm had also the best performance in two out of three measures.

Figure 4 shows that the ML algorithms performance had the same pattern in RQ1 and RQ2, as compared to RQ3.

D. RQ4: How ML predictions compare to user prediction?

We have compared ML algorithms predictions to user prediction in terms of error magnitude. Figure 5 shows predictors versus user error magnitude in the assignment of severity level. Figure 6 analyzes how well the ML prediction performed with respect to user prediction: better (ML algorithm error absolute value was smaller than user prediction error), equals (ML algorithm error equals to user error) or worse (ML algorithm error greater than user error). The data clearly show that the use of this type of software predictor results in no gain to the user. This conclusion could not be drawn simply knowing the value of the classic accuracy measurement for Neural Network (2426/4580 = 52.969%), Random Forest (2762/4580 = 60.305%), and SVM (2731/4580 = 59.628%) (see Figure 6). It is worth mentioning that our findings are in the same order of magnitude as findings reported in the literature. Therefore, one cannot state with confidence whether the use of the reported ML approach will bring any benefit, as compared to a simple educated guess by the user. On the contrary, there is evidence that predictions produced under these conditions are worse than user educated guess.

E. Statistical Tests

We can one observe in the previous tables that the performance of the three ML algorithms are very similar. To confirm whether they are really similar, we have evaluated their F-measure performances using Friedman Test [19]. We have defined the test null hypothesis (H0) as "the investigated algorithms have similar performance", and to test it, we have grouped F-measure values by dataset, research question and algorithm. For each group (3 per dataset), we have calculated the p-value. Table X shows the null hypothesis(H0) can be accepted (p$-$value $>= 0.05$) to the most cases, confirming that the investigated algorithms are similar [22].

TABLE X: Friedman tests results over F-measure.

	Question	P-value	H0
Cassandra	Q1	0.135335283	Accepted
	Q2	0.096971968	Accepted
	Q3	0.055637998	Accepted
Hadoop	Q1	0.223130160	Accepted
	Q2	0.096971968	Accepted
	Q3	0.006737947	Reject
Spark	Q1	0.223130160	Accepted
	Q2	0.096971968	Accepted
	Q3	0.040762204	Reject

F. Discussion

Table XI summarizes results related to CR severity level prediction reported in the literature and those generated during our experiments. We can observe that our results are better than others [4], [9] to address RQ1 and RQ2. we can also note that our results are better than [3] and worse than [8] to address RQ3. Notice that in many cases datasets and algorithms are different.

VI. CONCLUSIONS

In this paper, we have investigated the performance of three popular ML algorithms to predict CR severity level in an imbalanced data scenario. The results based on 22901 CRs extracted from the Cassandra, Hadoop e Spark repositories have shown that Random Forest results are slightly better than the other two algorithms to predict whether the severity level will change (RQ_1) and whether it will increase or decrease (RQ_2) with good F-measure, around 0.79 and 0.62 respectively, better than findings reported in the literature. On the other hand, results for the prediction of the final severity level on imbalanced data scenario (RQ_3) is similar to other results in the literature, with F-measure around 0.49. In an additional analysis conducted with Friedman Test, the three ML algorithms obtained similar performance for this experimental conditions. We have also shown that the classical measurements used in the literature do not help us deciding if the ML approach will bring any benefit to the user, and

Fig. 5: ML algorithms error magnitude (predictor versus user) for RQ3.

Fig. 6: ML algorithms performance compared to user for RQ4.

have proposed an alternative measuring approach to address this issue.

Validity threats to our research are: (a) We have assumed that user assigned severity level is correct and that there is

TABLE XI: ML algorithms performance summary (cells with a pair of numbers indicate range of variation).

	Research Question	Project	F-measure	Algorithm
Menzies [8]	Is the bug report blocker, critical, major, minor or trivial?	PitsA	14.0-17.0	Ripper
		PitsB	42.0-90.0	Ripper
		PitsC	53.0-92.0	Ripper
		PitsD	87.0-99.0	Ripper
		PitsE	8.0-88.0	Ripper
Lamkanfi [4]	Is the bug report severe or non-severe?	Mozilla	65.9-71.7	Naive Bayes
		Eclipse	62.5-65.5	Naive Bayes
		GNOME	72.7-78.5	Naive Bayes
Valdivia [9]	Is the bug report blocking or non-blocking?	Chromium	15.3	Decision Tree
		Eclipse	15.4	Decision Tree
		FreeDesktop	31.9	Decision Tree
		Mozilla	42.1	Decision Tree
		Netbeans	21.1	Decision Tree
		Netbeans	25.6	Decision Tree
Tian [3]	Is the bug report blocker, critical, major, minor or trivial?	OpenOffice	12.3-74.0	INSPect
		Mozilla	13.9-65.3	INSPect
		Eclipse	8.6-58.6	INSPect
Ours	Will the CR severity level change?	Cassandra	65.50-97.88	Random Forest
		Hadoop	53.16-97.09	Random Forest
		Spark	64.15-94.87	Random Forest
	Will the CR severity level increase, decrease or remain the same?	Cassandra	42.64-97.67	Random Forest
		Hadoop	18.51-97.26	Random Forest
		Spark	29.90-95.35	Random Forest
	Is the bug report blocker, critical, major, minor or trivial?	Cassandra	30.15-75.43	Random Forest
		Hadoop	30.76-82.01	Random Forest
		Spark	14.08-73.95	Random Forest

a close relationship between it and the long description of the CR. This assumption is supported [3], [4]. (b) We have considered three repositories and we have extracted 22901 CRs from it. Although we cannot generalize the results to others, the characteristics presented by Cassandra, Hadoop and Spark repositories, particularly regarding the balance of the data, are similar to those shown in the repositories studied [3]–[5], [9]. (c) Comparison in Table XI is relative, since it is based on experiments run on different datasets with different algorithms. (d) Code developed in Java language and the R language for preprocessing, training, testing and analysis of results have been carefully checked may still contain bugs.

As future work, we intend to investigate other repositories and systems, and develop an approach for representing CR Systems data in a general and uniform manner, so as to facilitate the development of a general purpose ML-based Maintenance Assistant.

REFERENCES

[1] Y. C. Cavalcanti, P. A. da Mota Silveira Neto, I. d. C. Machado, T. F. Vale, E. S. de Almeida, and S. R. d. L. Meira, "Challenges and opportunities for software change request repositories: a systematic mapping study," *Journal of Software: Evolution and Process*, vol. 26, no. 7, pp. 620–653, jul 2014.

[2] I. Sommerville, *Software Engineering*, 2010.

[3] Y. Tian, D. Lo, and C. Sun, "Information Retrieval Based Nearest Neighbor Classification for Fine-Grained Bug Severity Prediction," in *2012 19th Working Conference on Reverse Engineering*, oct 2012, pp. 215–224.

[4] A. Lamkanfi, S. Demeyer, E. Giger, and B. Goethals, "Predicting the severity of a reported bug," *Proceedings - International Conference on Software Engineering*, pp. 1–10, 2010.

[5] A. Lamkanfi, S. Demeyer, Q. D. Soetens, and T. Verdonckz, "Comparing mining algorithms for predicting the severity of a reported bug," *Proceedings of the European Conference on Software Maintenance and Reengineering, CSMR*, pp. 249–258, 2011.

[6] Y. Tian, D. Lo, X. Xia, and C. Sun, "Automated prediction of bug report priority using multi-factor analysis," *Empirical Software Engineering*, vol. 20, no. 5, pp. 1354–1383, oct 2015.

[7] N. V. Chawla, "Data Mining for Imbalanced Datasets: An Overview," in *Data Mining and Knowledge Discovery Handbook*. Boston, MA: Springer US, 2009, pp. 875–886.

[8] T. Menzies and A. Marcus, "Automated severity assessment of software defect reports," *IEEE International Conference on Software Maintenance, 2008. ICSM 2008*, pp. 346–355, 2008.

[9] H. Valdivia Garcia, E. Shihab, and H. V. Garcia, "Characterizing and predicting blocking bugs in open source projects," *Proceedings of the 11th Working Conference on Mining Software Repositories - MSR 2014*, pp. 72–81, 2014.

[10] R. S. Pressman, *Software Engineering A Practitioner's Approach 7th Ed - Roger S. Pressman*, 2009.

[11] R. Feldman and J. Sanger, *The Text Mining Handbook: Advanced Approaches in Analyzing Unstructured Data*. Cambridge University Press, 2007.

[12] G. Williams, *Data Mining with Rattle and R: The Art of Excavating Data for Knowledge Discovery*, 2011. [Online]. Available: http://books.google.com/books?id=mDs7OXj03V0C

[13] K. Surya, R. Nithin, and R. Venkatesan, "A Comprehensive Study on Machine Learning Concepts for Text Mining," *International Conference on Circuit, Power and Computing Technologies [ICCPCT] A*, vol. 3, no. 1, pp. 1–5, 2016.

[14] S. J. Russell and P. Norvig, *Artificial Intelligence: A Modern Approach*, 2010.

[15] L. Breiman, "Random Forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.

[16] K. Facelli, A. C. Lorena, J. Gama, and A. Carvallho, *Inteligência Artifical: uma abordagem de aprendizado de máquina*. Rio de Janeiro: LTC, 2015.

[17] Y. Zhao and Y. Cen, *Data Mining Applications with R*, 1st ed. Academic Press, 2013.

[18] E. de Jonge and M. van der Loo, "An introduction to data cleaning with R," *Statistics Netherlands*, p. 53, 2013. [Online]. Available: http://cran.r-project.org/doc/contrib/de_Jonge+van_der_Loo-Introduction_to_data_cleaning_with_R.pdf

[19] N. Japkowicz and M. Shah, *Evaluating Learning Algorithms: A Classification Perspective*. New York, NY, USA: Cambridge University Press, 2011.

[20] S. Haykin, *Neural Networks: A Comprehensive Foundation*, 2nd ed. Upper Saddle River, NJ, USA: Prentice Hall PTR, 1998.

[21] N. Cristianini and J. Shawe-Taylor, *An Introduction to Support Vector Machines: And Other Kernel-based Learning Methods*. New York, NY, USA: Cambridge University Press, 2000.

[22] J. Demšar, "Statistical comparisons of classifiers over multiple data sets," *J. Mach. Learn. Res.*, vol. 7, pp. 1–30, Dec. 2006. [Online]. Available: http://dl.acm.org/citation.cfm?id=1248547.1248548